



ShadowIQ
NAVIGATE THE UNKNOWN

OUR METHODOLOGY

How We Score Risk

Four separate questions — how bad could it be, is it still dangerous, are your people in its way, and how sure are we? Keeping them apart is what stops one confident-sounding number from hiding what matters.

THE SHADOWIQ METHODOLOGY IN PLAIN LANGUAGE



Four questions, kept apart on purpose

The common failure in risk scoring is to mash everything into one number — so a confirmed earthquake that finished yesterday looks the same as a wildfire heading your way, and a rock-solid government advisory looks the same as an anonymous social post. ShadowIQ scores each question on its own, then brings them together.

CONSEQUENCE

1

How bad could it be?

The worst realistic harm to a person — from minor disruption to life-threatening.

LIKELIHOOD

2

Is it still dangerous?

How likely the hazard is to affect someone from now on — not whether it already happened.

EXPOSURE

3

Are your people in its way?

Whether a trip is scheduled in the affected place, during the affected time window.

CONFIDENCE

4

How sure are we?


How much we trust the reporting — kept strictly apart from how risky the event is.

1

Consequence — how bad could it be?

If this hazard reached one of your travellers, how bad could it realistically be? ShadowIQ places every event on a plain severity ladder — from negligible up to catastrophic — anchored to the language government travel advisories already use, so grading stays consistent from one event to the next.

Consequence: harm to a person, on a 1–5 ladder




2

Likelihood — is it still dangerous?

Likelihood is **not** “did this happen?” Almost everything we ingest genuinely happened. The question that matters is: *is there still a hazard ahead of the traveller?* An earthquake that struck yesterday with no ongoing danger is history; a wildfire still spreading is not. We describe it with a defined vocabulary drawn from **ICD 203**.


One shared vocabulary: the **ICD 203** probability yardstick



3

Exposure — are your people in its way?


A live hazard on the other side of the world is not your problem; the same hazard where your group is standing very much is. ShadowIQ does **not** track travellers' phones — it matches events to the **places and dates** your trips are scheduled. An event only counts if a trip is in the affected place during the affected time window.



4

Confidence — how sure are we?


A government advisory and an anonymous social post are not equal evidence. ShadowIQ grades every signal with the **NATO Admiralty system** — source reliability and claim corroboration, scored independently. Crucially, confidence is kept **separate** from risk: a shaky report of something serious is routed to a human analyst to verify, not inflated and not buried.






How it comes together into a trip level

A trip's overall level is set by the most serious hazard that is *currently live and actually relevant* to that itinerary — a **high-water mark**, not an average that dilutes one severe threat among many mild ones. And because likelihood eases as hazards resolve, the level comes back down on its own once the danger has passed.



AI and humans keep it honest

AI does the first pass at machine scale. But “AI-powered” is easy to say and hard to trust, so every serious alert is reviewed by a human — the automation cannot close those on its own — and a continuous, blind-sampled slice of ordinary alerts is human-checked too. Those decisions are the reference standard the scoring is tuned against, and the agreement between AI and analysts is **measured, not assumed**.





Who gets told, and how fast

Alert fatigue is a safety failure. ShadowIQ notifies immediately only for **credible alerts that intersect an itinerary**; uncertain-but-serious reports get a **human check before they interrupt anyone**; the rest arrives in digests.

Quiet is a feature — it is what makes the loud moments trustworthy. A tool that interrupts people constantly trains them to ignore it, and the one alert that mattered goes unread.



Built on standards, not invented from scratch

The bones of the model come from references your auditors, insurers, and security leads already recognise:

ISO 31030

Travel risk management — likelihood × consequence, assessed for the individual traveller on a specific itinerary.


ICD 203

The intelligence community's analytic standard — a defined probability vocabulary; never blend likelihood with confidence.

NATO Admiralty

Source reliability and claim corroboration, graded independently. Defensible beats opaque.

Standards are a floor, not a ceiling. The value we add is applying them at machine scale, continuously — and calibrating that automation against human analyst judgement.



Where our job ends

ShadowIQ is decision-support: it helps you *see* and *understand* risk so the right people can act with clear, defensible evidence. It is not a live-location tracker, and it does not run the response — we do not coordinate evacuations, manage the incident, or make the duty-of-care decision on your behalf. Those calls stay where the accountability sits: with you.



ShadowIQ
NAVIGATE THE UNKNOWN

Read the full [Inside the Risk Model](#) series · shadowiq.io/blog · shadowiq.io/how-we-score-risk · hello@shadowiq.io